

Counting Triangles in Large Graphs using Randomized Matrix Trace Estimation: an Overview

Steven Haussmann, Louis Hyde, Peter Wood

Dec. 3, 2018

Background

A **triangle** is a group of three nodes in a graph that are pairwise connected

- ▶ Alice, Bob, and Charlie are all friends with one another
- ▶ Three cities are connected by freeways

- ▶ The number of triangles in a graph provides interesting information:
 - ▶ Clustering coefficient - how likely are nodes to be linked?
- ▶ Iterating over nodes or edges takes $O(\sum_{v \in V} \deg(v)^2) \approx O(|E|^2)$ time and is hard to parallelize
- ▶ Monte Carlo simulation to check if random nodes/edges form a triangle only works for dense graphs

We examined a 2010 paper on the subject, *Counting Triangles in Large Graphs using Randomized Matrix Trace Estimation*

- ▶ Authored by Haim Avron of Tel-Aviv University

Avron presents a *randomized approach* to triangle counting.
We examine two of his proposed algorithms.

Algorithm

Randomized Trace Estimation

- ▶ A is the adjacency matrix
- ▶ x is a vector whose entries are independent random variables with expectation 0 and variance 1
- ▶ Calculate $y = Ax$
- ▶ Then $y^T Ay = (x^T A)A(Ax) = x^T A^3 x$ approximates $\text{tr}(A^3)$

Randomized Trace Estimation, cont'd

- ▶ $\mathbb{E}(x^T A x) = \sum_{i=0}^n \sum_{j=0}^n \mathbb{E}(x_i A_{ij} x_j)$
- ▶ Since $\mathbb{E}x_i = 0$, $\mathbb{E}(x_i x_j) = 0$ for $i \neq j$, so $\sum_{i \neq j} \mathbb{E}(x_i A_{ij} x_j) = 0$
- ▶ Since $\mathbb{E}(x_i^2) = \mathbb{E}(x_i^2) + \text{Var}(x_i) = 1$, $\mathbb{E}(x_i A_{ii} x_i) = A_{ii}$
- ▶ Thus $\mathbb{E}(x^T A x) = \sum_{i=0}^n A_{ii} = \text{Tr}(A)$

Two variants of the same algorithm were presented by the author:

- ▶ TraceTriangleN (normal RVs)
- ▶ TraceTriangleR (Rademacher RVs)

TraceTriangleN

Result: Δ , the estimated number of triangles

Form an adjacency matrix $A \in R^{n \times n}$;

$M \leftarrow \lceil \gamma \ln^2 n \rceil$;

$i \leftarrow 0$;

while $i < M$ **do**

$x \leftarrow (x_k)$ where $x_i \sim N(0, 1)$ (x_i are i.i.d.);

$y \leftarrow Ax$;

$T_i \leftarrow (y^T Ay)/6$;

$i \leftarrow i + 1$;

end

$\Delta \leftarrow \frac{1}{M} \sum_{i=1}^M T_i$

Variants of the algorithm:

- ▶ Multiple ways to sample x
- ▶ Author notes that $N(0, 1)$ gives the best bound
- ▶ Rademacher approach is better for implementation, however
- ▶ Also theoretically has a better estimator of the trace

In summary:

- ▶ Compute the average of M samples
- ▶ Samples can be taken in parallel
- ▶ Experimental results suggest $M = \lceil \gamma \ln^2 |V| \rceil$

Analysis

Speedup

- ▶ Any exact algorithm has to involve matrix multiplication, which is naïvely $O(|V|^3)$
- ▶ Better multiplication algorithms live somewhere above $O(|V|^2)$
- ▶ TraceTriangle requires $O(|E|)$ time for each sample
 - ▶ Requires $A \times x$ for each iteration, which would make it $O(|V|^2)$ but if A is sparse, this could be made $O(|E|)$
- ▶ $\lceil \gamma \log^2(|V|) \rceil$ samples are needed
- ▶ Total runtime: $O(\gamma |E| \log^2 |V|)$

Avron describes (ε, δ) bounds on error.

- ▶ Depends on $\rho(G)$ - the *eigenvalue triangle sparsity* of the graph.
- ▶ $\rho(G) = \frac{\sum_{i=1} n |\lambda_i|^3}{6\Delta(G)}$
- ▶ Very small triangle counts and very large eigenvalues cause large error bounds
- ▶ γ required for $\varepsilon = 0.05$, $\delta = 0.05$ is enormous
- ▶ Author notes that small γ works well *in practice* for commonly encountered graphs.

Random Bits

- ▶ TraceTriangleN requires n random numbers per step
- ▶ $O(\gamma \log^2 |V|)$ total bits
- ▶ TraceTriangleR is similar, requiring n random bits
- ▶ Avron proposed a third algorithm requiring significantly fewer; not examined here

Random number generation may seem trivial, but becomes difficult in massively parallel situations

Scalability

TraceTriangleN is *embarrassingly parallel*

- ▶ Each iteration is independent of all other iterations
- ▶ Can easily farm out to many processors

The algorithm can also be streamed.

- ▶ One pass requires $O(|V| \log^2 |V|)$ memory
- ▶ Less memory is consumed if several passes are made

We evaluated the performance of TraceTriangleN and TraceTriangleR in three situations:

- ▶ Runtime on graphs of varying size
- ▶ Runtime and variance of results with varying γ
- ▶ Variance and error of results with varying triangle sparsity

Notes:

- ▶ The exact algorithm directly computes $\text{Tr}(A^3)$
- ▶ `TraceTriangleN` and `TraceTriangleR` are not parallelized
- ▶ $\gamma = 1$ unless otherwise noted
- ▶ Sparse matrix multiplication was not performed (increasing runtime)

Runtime

We evaluated runtime for graphs of varying size.

- ▶ $|V| = 1000, 2000, \dots, 10000$
- ▶ $\gamma = 1$
- ▶ 25 trials

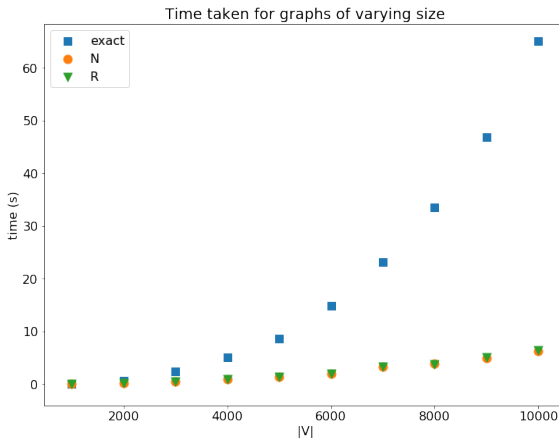


Figure: Exactly computing $\text{Tr}(A^3)$ proves to be far more expensive than the randomized approach.




Figure: $\gamma = 1$, but values presented are the average of 25 trials (similar to letting $\gamma = 25$). We examine accuracy more closely next.

Gamma

We examined the variance and average absolute percent error, where:

- ▶ $|V| = 1000$
- ▶ A is tridiagonal
- ▶ $\gamma = 1, 2, \dots, 10$
- ▶ 100 trials were performed for each value of γ

Empirical Results




Figure: As expected, time taken scales linearly with γ




Figure: Variance declines as γ increases. Note the apparent noise, even with 100 trials.

Empirical Results




Figure: Average absolute percent error also declines.

Sparsity

We examined the variance and average absolute percent error, where:

- ▶ $|V| = 100$
- ▶ $\max(\Delta(G)) = 100 \times 99 \times 98/6 = 161700$
- ▶ A has n number of triangles added randomly:
 - ▶ Select random distinct i, j, k
 - ▶ Connect vertices i, j, k together
- ▶ $n = [2^0, 2^{\frac{1}{2}}, 2^1, \dots, 2^{20}]$
- ▶ $\gamma = 1$
- ▶ 25 trials were performed for each value of n




Figure: Note how each added triangle may actually create more than one triangle, and how the graph becomes “saturated”




Figure: The variance grows with the number of triangles, as expected. A slight dip is observed in the sparser regions.

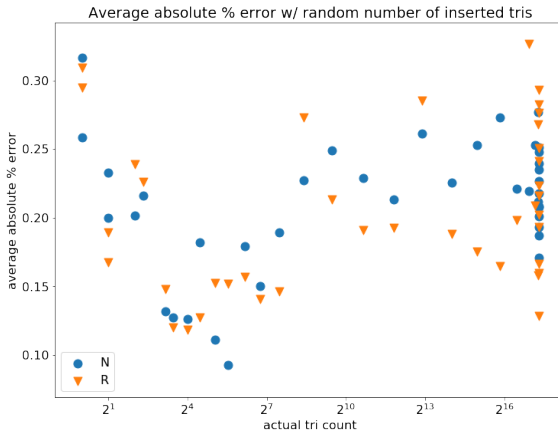


Figure: The average absolute percent error is lowest when the graph is sparse (but not *too* sparse). Note that the repeated trials for a completely full graph gave widely varying results.

Conclusions

Objective

- ▶ TraceTriangle offers significant speedup over the naïve approach
- ▶ Parallelization is straightforward, although bandwidth is a concern.
- ▶ Accuracy is reasonable
 - ▶ Being an unbiased estimator means many trials can be averaged
- ▶ Somewhat questionable theoretical basis - low values of γ don't provide strong guarantees, but do work on real data.

Subjective

- ▶ *Counting Triangles* has over 50 citations
- ▶ Builds on both deterministic and randomized work
- ▶ 2010 paper, so work has been done since publication